

# Advanced Design Patterns for Amazon DynamoDB

Sean Shriver, AWS DynamoDB SA  
April, 2020

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



1

## What to Expect from the Session?

- DynamoDB 101
- Tenants of NoSQL Data Modeling
  - Normalized versus Denormalized schemas
- Common Design Patterns
  - Working Backwards: Building Queries
  - Indexes, Filters, and Hierarchies
  - Write Sharding, Distributed Transactions, & Complex Queries
- DynamoDB in the AWS Ecosystem

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



2

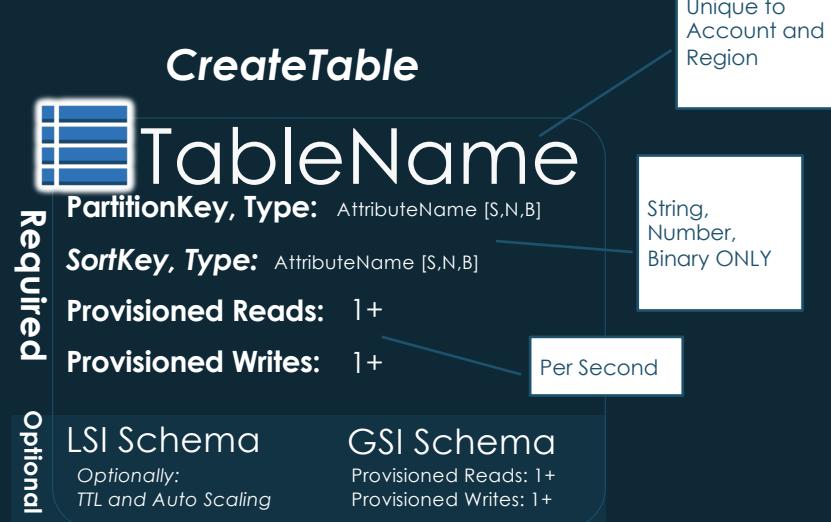
# DynamoDB 101

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



3

## Table creation options



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



4

# Provisioned capacity

## Provisioned capacity

*Capacity is per second, rounded up to the next whole number*

### Read Capacity Unit (RCU)

1 RCU returns 4KB of data for strongly consistent reads, or double the data at the same cost for eventually consistent reads

### Write Capacity Unit (WCU)

1 WCU writes 1KB of data, and each item consumes 1 WCU minimum

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



5

# On demand capacity

Read/write capacity on demand



### No capacity planning

No need to specify how much read/write throughput you expect to use



### Ideal for unpredictable workloads

Ramp from zero to tens of thousands of requests per second on demand



### Pay only for what you use

Pay-per-request pricing

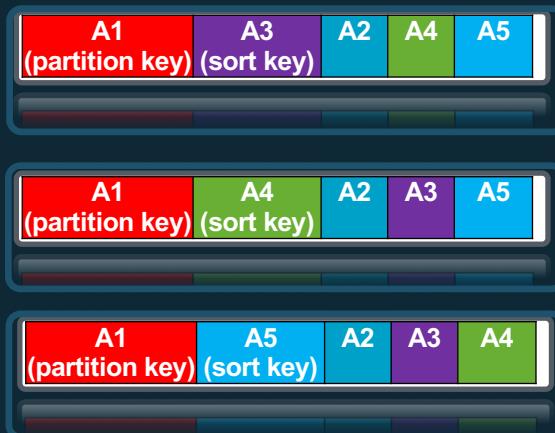
© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



6

## Local Secondary Indexes

- Alternate sort key attribute
- Index is local to a partition key



10 GB max per partition key, i.e.  
LSIs limit the # of sort keys!

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



7

## Global Secondary Indexes

- Alternate partition (+sort) key
- Index is across all table partition keys
- Can be added or removed anytime



RCUs/WCUs provisioned separately for GSIs

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



8

# Horizontal Sharding

## CustomerOrdersTable



**~Each new host brings compute, storage and network bandwidth~**

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



9

# Partitioning

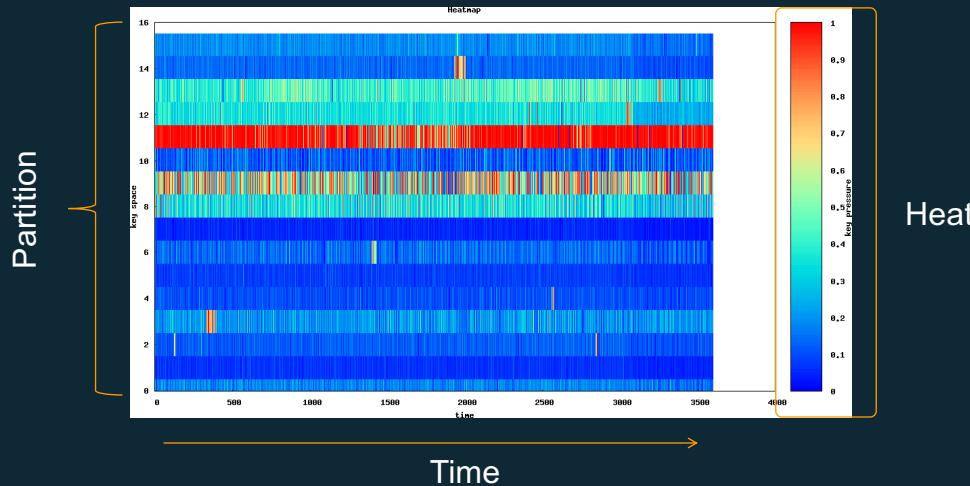


© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



10

## What bad NoSQL looks like...



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



11

## Getting the most out of DynamoDB throughput

"To get the most out of DynamoDB throughput, create tables where the partition key element has a large number of distinct values, and values are requested fairly uniformly, as randomly as possible."

—*DynamoDB Developer Guide*

---

**Space:** access is evenly spread over the keyspace

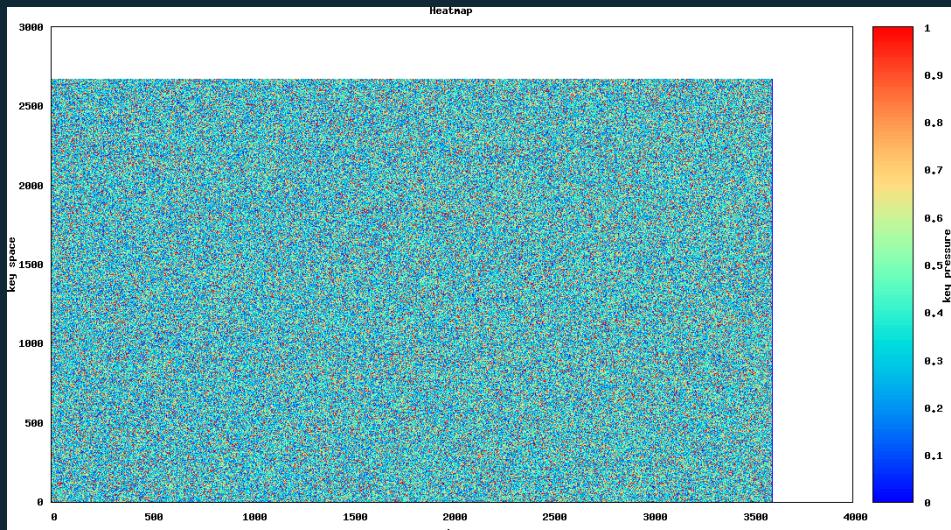
**Time:** requests arrive evenly spaced in time

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



12

Much better picture...



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



13

## Amazon DynamoDB

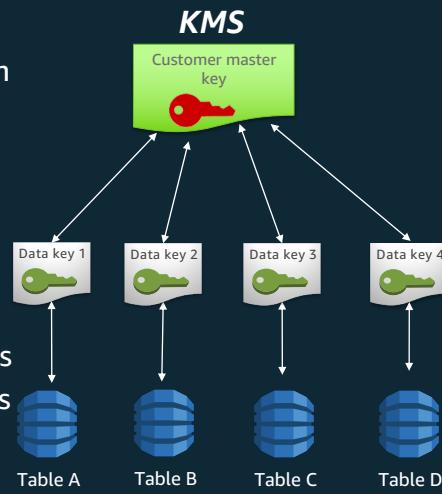
### Encryption at rest

Two-tiered key hierarchy using envelope encryption

- Unique data key encrypts customer data
- Customer Master Keys encrypt data keys

#### Benefits

- Encrypts both base tables and indexes
- Transparent, no application changes
- Better performance for encrypting large tables
- Easier to manage small number of master keys than billions of data keys
- Centralized access and audit of key activity



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



14

## Amazon DynamoDB – Backup and Restore

Only cloud database to provide on demand and continuous backups



On-demand  
backups for long-  
term data archival  
and compliance



Point in time restore  
for short term  
retention and data  
corruption protection  
(35 days)



Point in time recovery with  
restore times in a few hours  
depending on table size

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



15

## Adaptive Capacity

- Dynamically adjusts partition capacity based on real-time traffic
  - To better handle imbalanced workloads
  - Best practices for even load distribution still matter
- Reduces throttling
  - As long as there is enough capacity provisioned for the table
  - Up to the hard limit of partition write capacity
  - Triggered by throttling

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



16

# DynamoDB Global Tables

Fully managed, multi-master, multi-region database



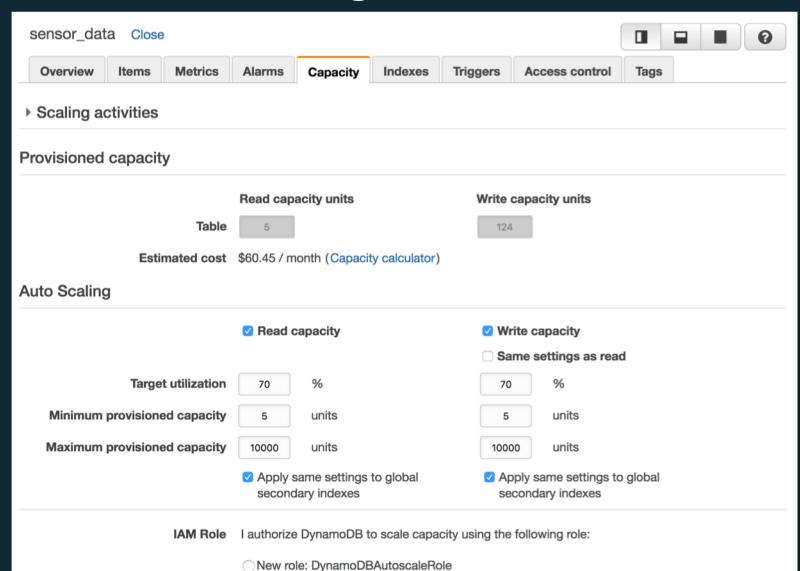
- Build high performance, globally distributed applications
- Low latency reads & writes to locally available tables
- Disaster proof with multi-region redundancy
- Easy to setup and no application re-writes required

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

**aws**

17

## DynamoDB Auto Scaling



sensor\_data Close

Overview Items Metrics Alarms Capacity Indexes Triggers Access control Tags

Scaling activities

Provisioned capacity

Table	Read capacity units	Write capacity units
sensor_data	5	124

Estimated cost \$60.45 / month ([Capacity calculator](#))

Auto Scaling

<input checked="" type="checkbox"/> Read capacity	<input checked="" type="checkbox"/> Write capacity
<input type="checkbox"/> Same settings as read	<input type="checkbox"/> Same settings as read
Target utilization 70 %	70 %
Minimum provisioned capacity 5 units	5 units
Maximum provisioned capacity 10000 units	10000 units
<input checked="" type="checkbox"/> Apply same settings to global secondary indexes	<input checked="" type="checkbox"/> Apply same settings to global secondary indexes

IAM Role I authorize DynamoDB to scale capacity using the following role:

New role: DynamoDBAutoscaleRole  
 Existing role with pre-defined policies [\[Instructions\]](#)

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

**aws**

18

# NoSQL data modeling

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



19

It's all about aggregations...



Social network



Document management



Process control



IT monitoring

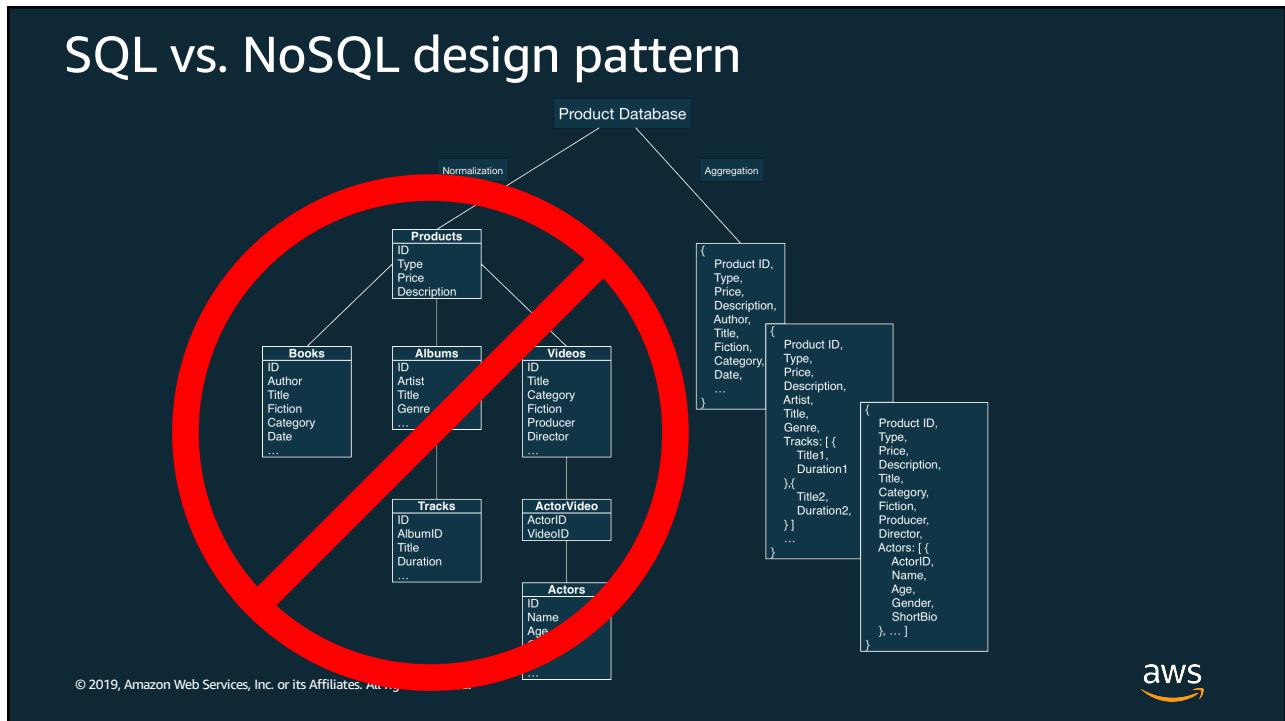


Data trees

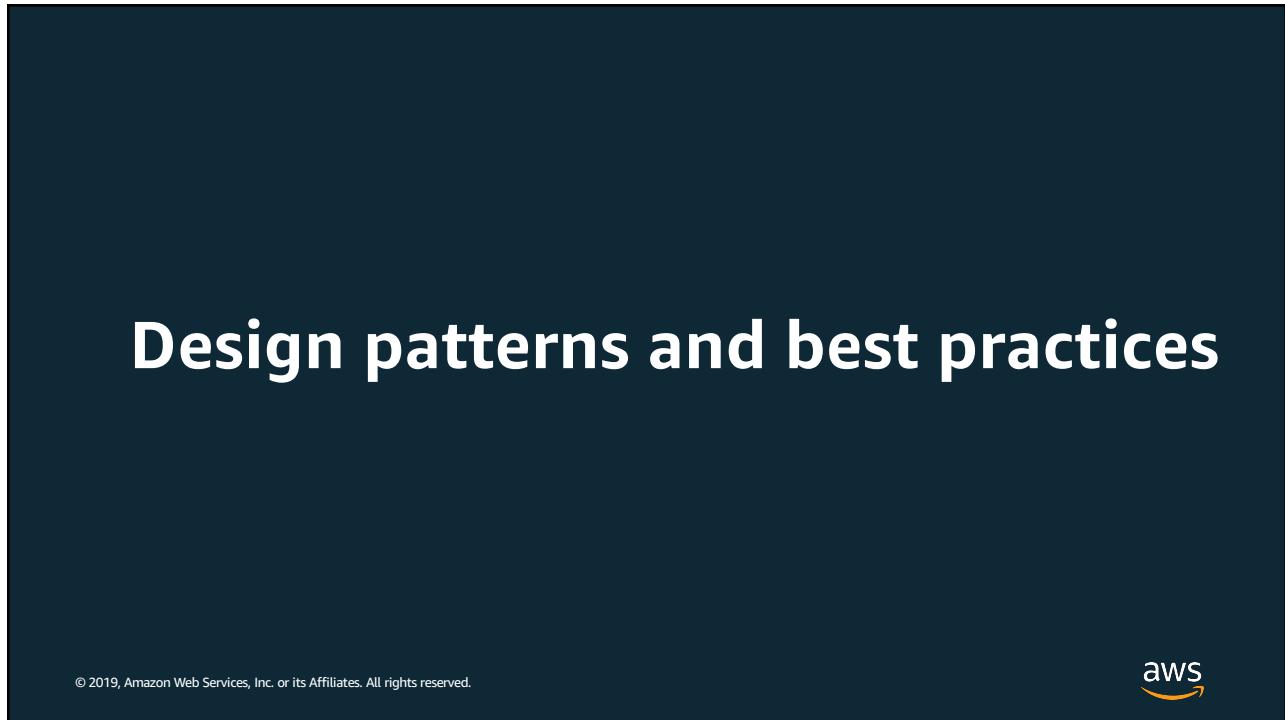
© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



20



21



22

## TENETS OF NoSQL DATA MODELING

- Understand the use case
- Identify the access patterns
  - Read/Write workloads
  - Query dimensions and aggregations
- Data-modeling
  - Avoid relational design patterns, use one table
- Review -> Repeat -> Review

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



23

## TENETS OF NoSQL DATA MODELING

- Understand the use case
  - Identify the access patterns
    - Read/Write workloads
    - Query dimensions and aggregations
  - Data-modeling
    - Avoid relational design patterns, use one table
  - Review -> Repeat -> Review
- Nature of the application
    - OLTP / OLAP
  - Define the Entity-Relationship Model
  - Identify Data Life Cycle
    - TTL, Backup/Archival, etc.

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



24

## TENETS OF NoSQL DATA MODELING

- Understand the use case
- Identify the access patterns
  - Read/Write workloads
  - Query dimensions and aggregations
- Data-modeling
  - Avoid relational design patterns, use one table
- Review -> Repeat -> Review
- Identify data sources
- Define queries and write patterns
- Document all workflows

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



25

## TENETS OF NoSQL DATA MODELING

- Understand the use case
- Identify the access patterns
  - Read/Write workloads
  - Query dimensions and aggregations
- Data-modeling
  - Avoid relational design patterns, use one table
- Review -> Repeat -> Review
- **1 application service = 1 table**
  - Reduce round trips
  - Simplify access patterns
- Identify Primary Keys
  - How will items be inserted and read?
  - Overload items into partitions
- Define indexes for secondary access patterns

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



26

# Building Queries

Query filters, composite keys, and sparse indexes

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



27

## Multi-value sorts and filters

**Partition key**

**Sort key**

Base table index

GamerID	Date	Gameld	Status	Host
Alice	2018-10-02	d9bl3	DONE	David
Carol	2018-10-08	o2pnb	IN_PROGRESS	Bob
Bob	2018-09-30	72f49	PENDING	Alice
Bob	2018-10-03	b932s	PENDING	Carol
Bob	2018-10-03	ef9ca	IN_PROGRESS	David



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



28

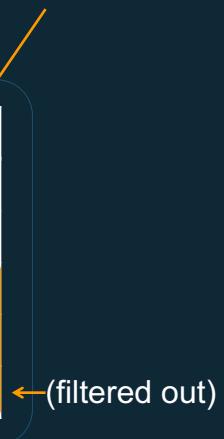
## Approach 1: Query filter

```
SELECT * FROM Game
WHERE Opponent='Bob'
ORDER BY Date DESC
FILTER ON Status='PENDING'
```



Base table index

GamerID	Date	GameId	Status	Host
Alice	2018-10-02	d9bl3	DONE	David
Carol	2018-10-08	o2pnb	IN_PROGRESS	Bob
Bob	2018-09-30	72f49	PENDING	Alice
Bob	2018-10-03	b932s	PENDING	Carol
Bob	2018-10-03	ef9ca	IN_PROGRESS	David



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



29

## Approach 2: Composite key

Status	Date	StatusDate
DONE	2018-10-02	DONE_2018-10-02
IN_PROGRESS	2018-10-08	IN_PROGRESS_2018-10-08
IN_PROGRESS	2018-10-03	IN_PROGRESS_2018-10-03
PENDING	2018-10-03	PENDING_2018-10-03
PENDING	2018-09-30	PENDING_2018-09-30

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



30

## Approach 2: Composite key

Partition key      Sort key

 Base table index

GamerID	StatusDate	Gameld	Host
Alice	DONE_2018-10-02	d9bl3	David
Carol	IN_PROGRESS_2018-10-08	o2pnb	Bob
Bob	IN_PROGRESS_2018-10-03	ef9ca	David
Bob	PENDING_2018-09-30	72f49	Alice
Bob	PENDING_2018-10-03	b932s	Carol

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



31

## Approach 2: Composite key

```
SELECT * FROM Game
WHERE Opponent='Bob'
AND StatusDate BEGINS WITH 'PENDING'
```

 Base table index

GamerID	StatusDate	Gameld	Host
Alice	DONE_2018-10-02	d9bl3	David
Carol	IN_PROGRESS_2018-10-08	o2pnb	Bob
Bob	IN_PROGRESS_2018-10-03	ef9ca	David
Bob	PENDING_2018-09-30	72f49	Alice
Bob	PENDING_2018-10-03	b932s	Carol



Bob

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



32

## Design for the access patterns:

```
SELECT * FROM CartTable
WHERE CartID = 'BobCT'
    AND StatusDate BEGINS_WITH 'ACTIVE'
```

 Bob's cart

Partition Key	Sort Key	ItemSKU	Count	UnitPrice	Gift
CartID	Status Date				
AliceCart	SAVED_2018-10-02	d9bl3		34.99	
CarolCart	SHIPPED_2018-10-08	o2pnB		16.99	
BobCart	SAVED_2018-10-03	ef9ca		4.99	
BobCart	ACTIVE_2018-09-30	72f49	2	12.49	
BobCart	ACTIVE_2018-10-03	b932s	1	67.99	Alice

GSI SortKey      GSI Partition Key

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

33

## Sparse indexes

Game-scores-table					
Id (Partition)	User	Game	Score	Date	Award
1	Bob	G1	1300	2012-12-23	
2	Bob	G1	1450	2012-12-23	
3	Jay	G1	1600	2012-12-24	
4	Mary	G1	2000	2012-10-24	Champ
5	Ryan	G2	123	2012-03-10	
6	Jones	G2	345	2012-03-20	

Scan sparse GSIs

Award-GSI			
Award (Partition)	Id	User	Score
Champ	4	Mary	2000

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

34



## Replace filter with indexes



Concatenate attributes to form useful secondary index keys



Take advantage of sparse indexes

**Important when:** you want to optimize a query as much as possible

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



35

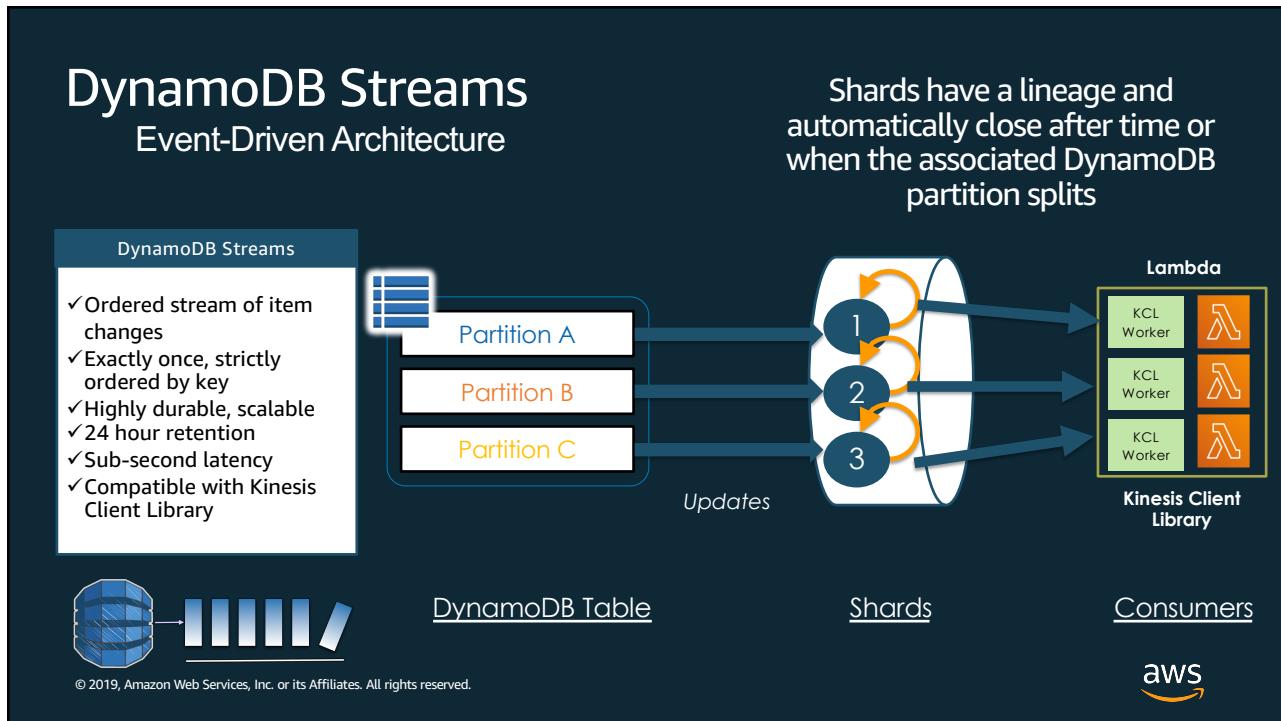
## Complex Queries & Write Sharding

Realtime aggregations  
Rebalancing write distribution

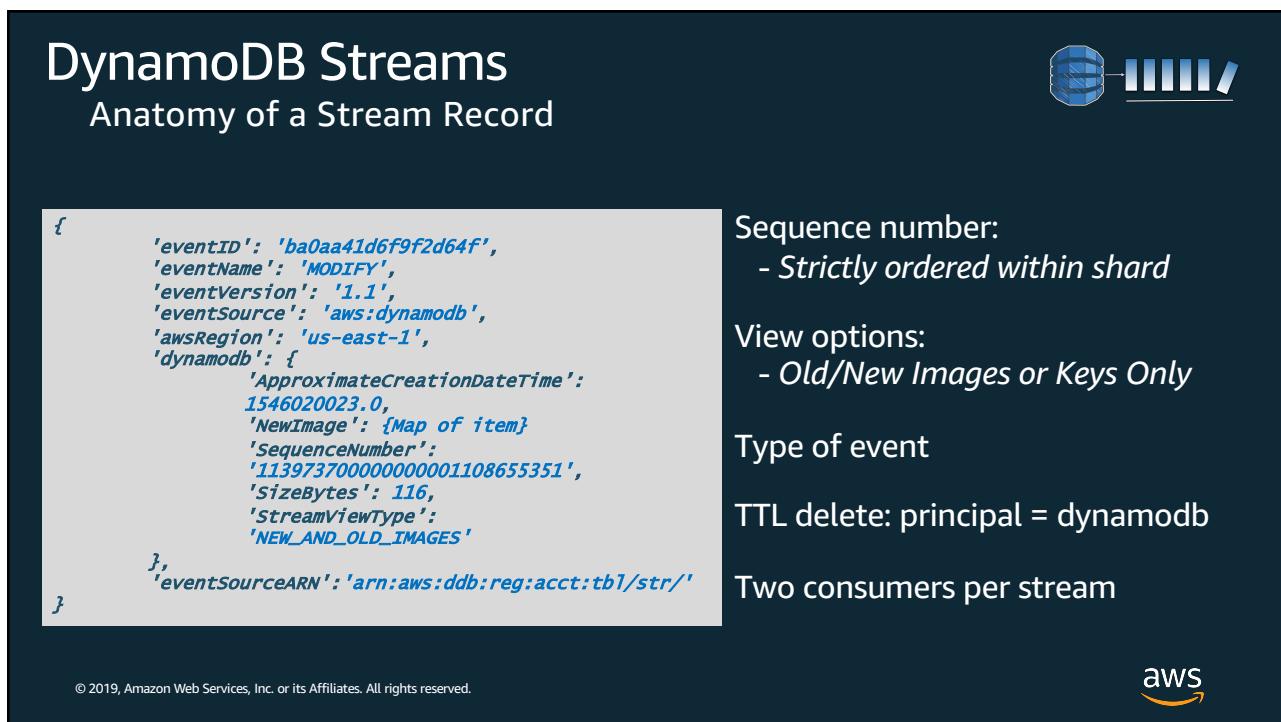
© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



36

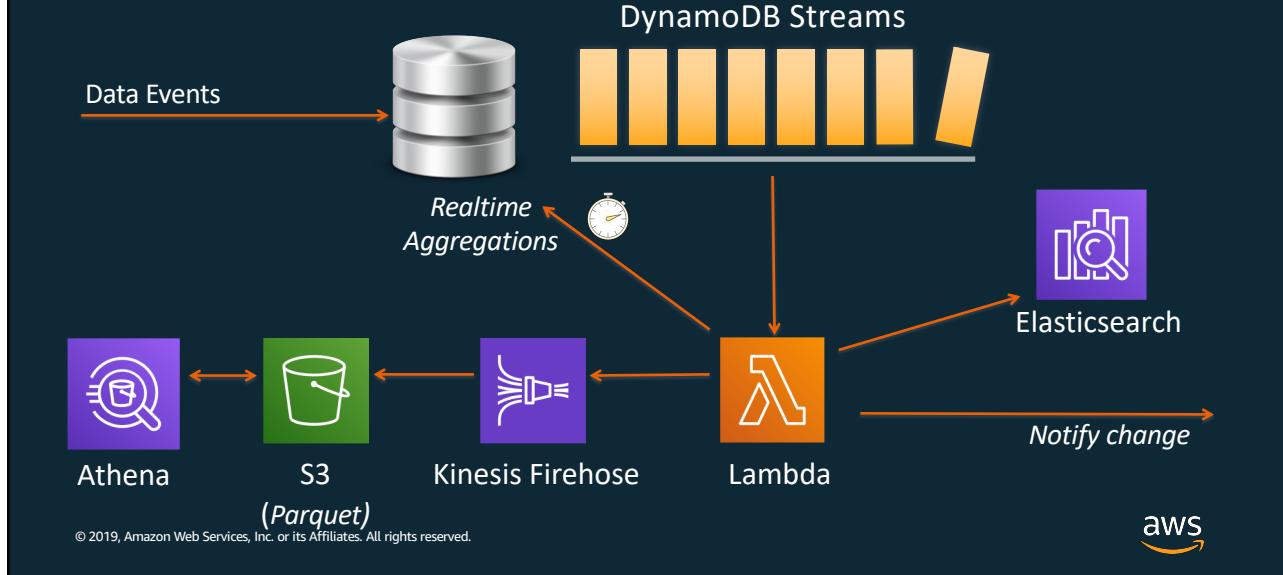


37

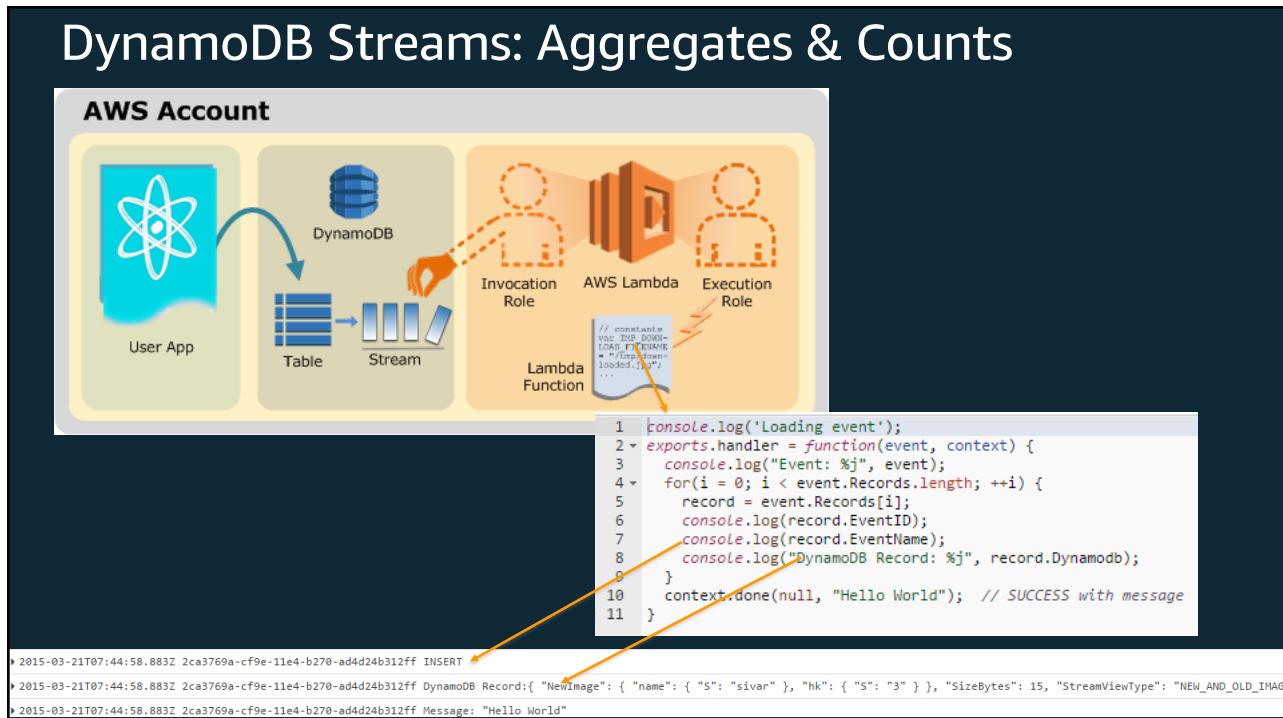


38

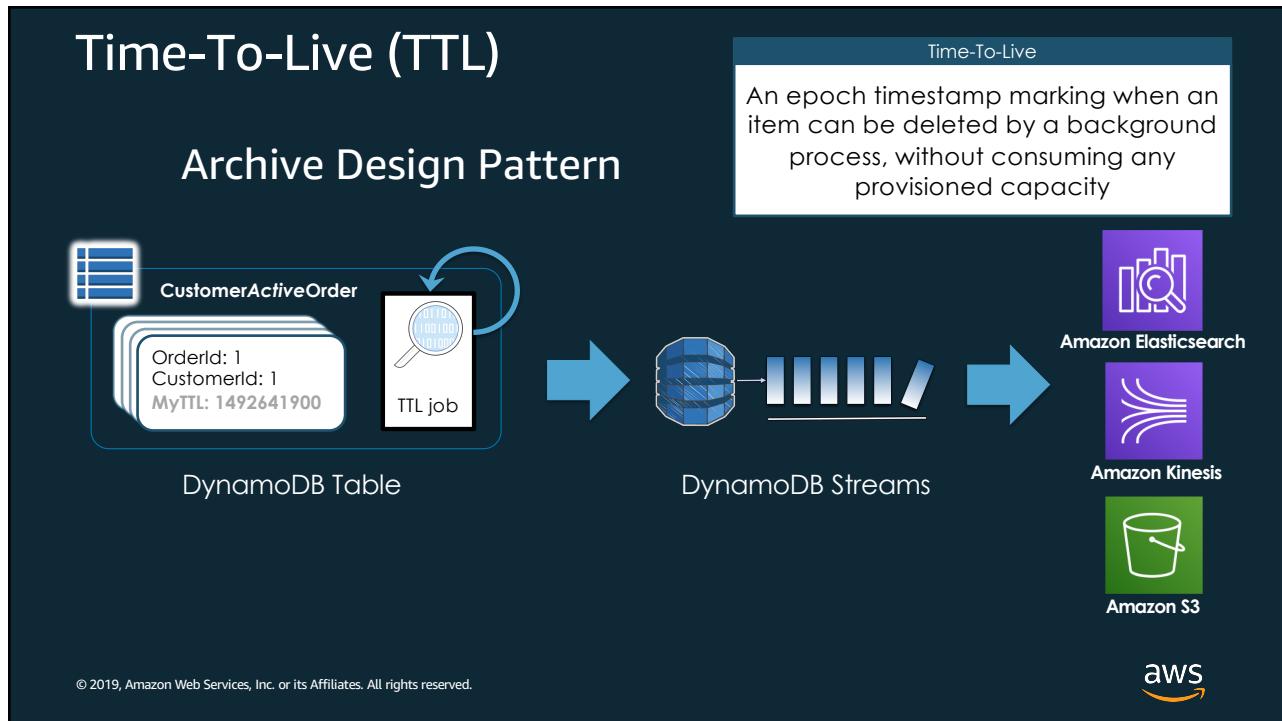
# Serverless & Event Driven Architecture



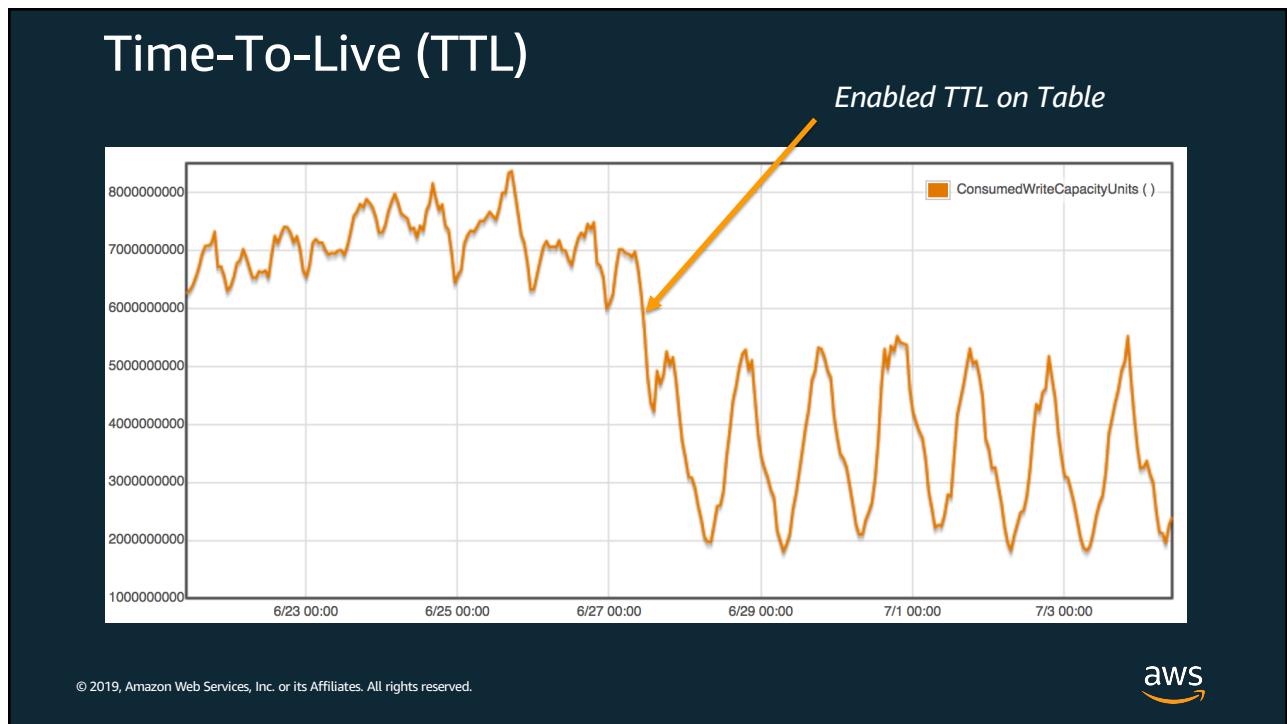
39



40



41



42

# Write sharding

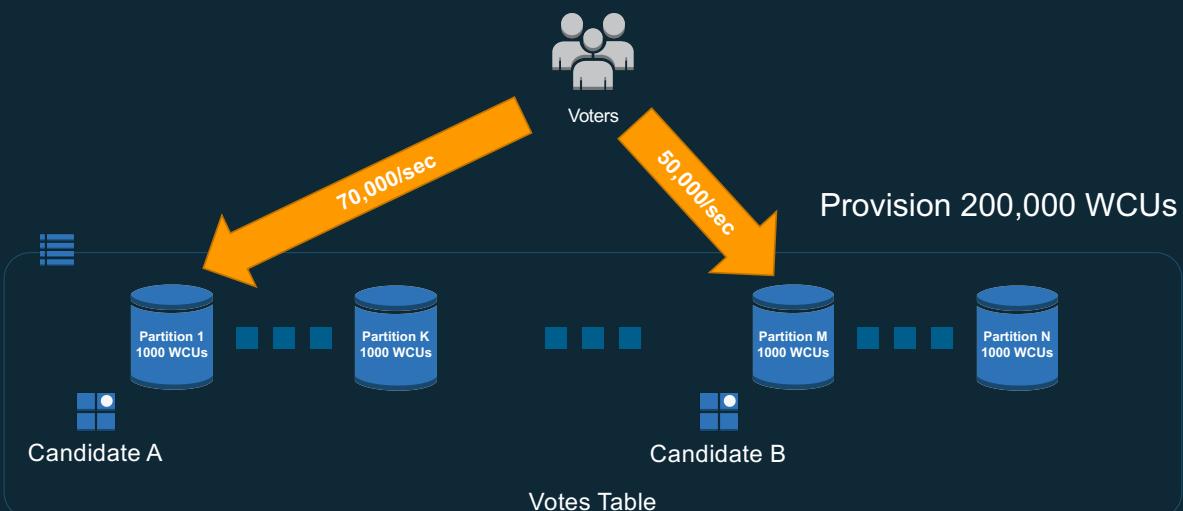
Handling writes with an access distribution skew

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



43

## Scaling bottlenecks



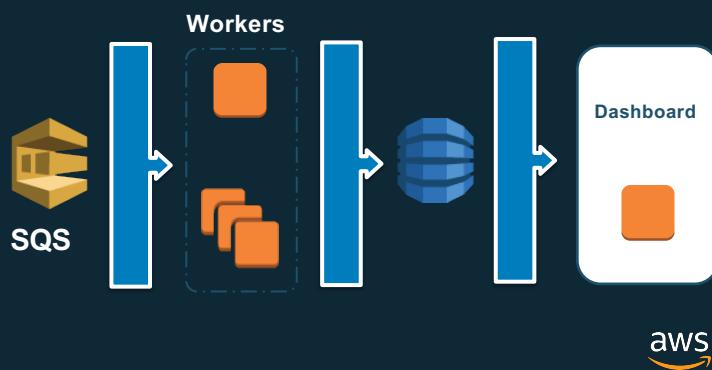
© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



44

## Queue-based load leveling

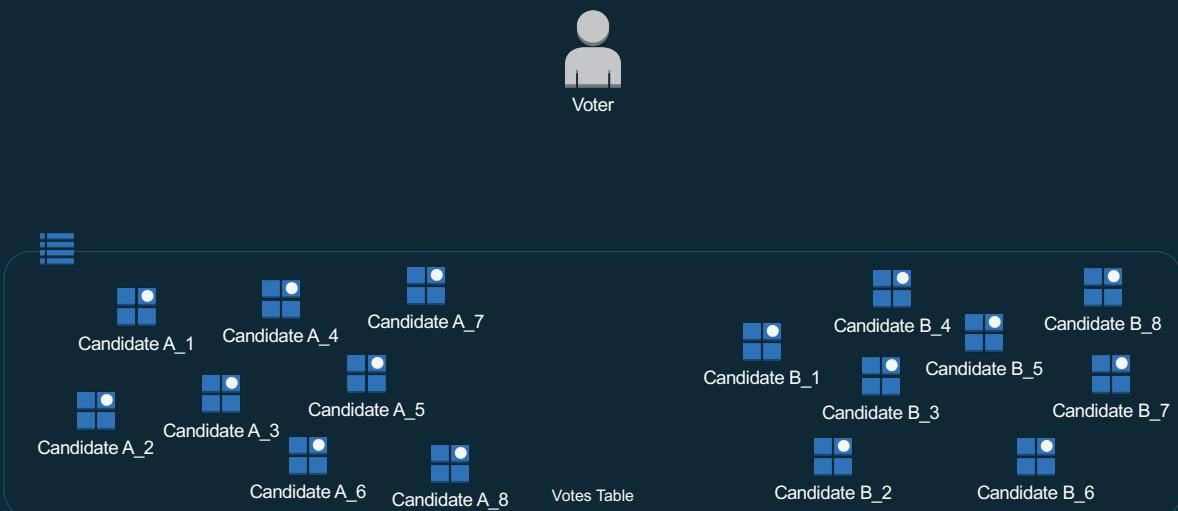
- Write incoming votes to SQS
- Scale workers as needed to process queue
- Avoid key pressure



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

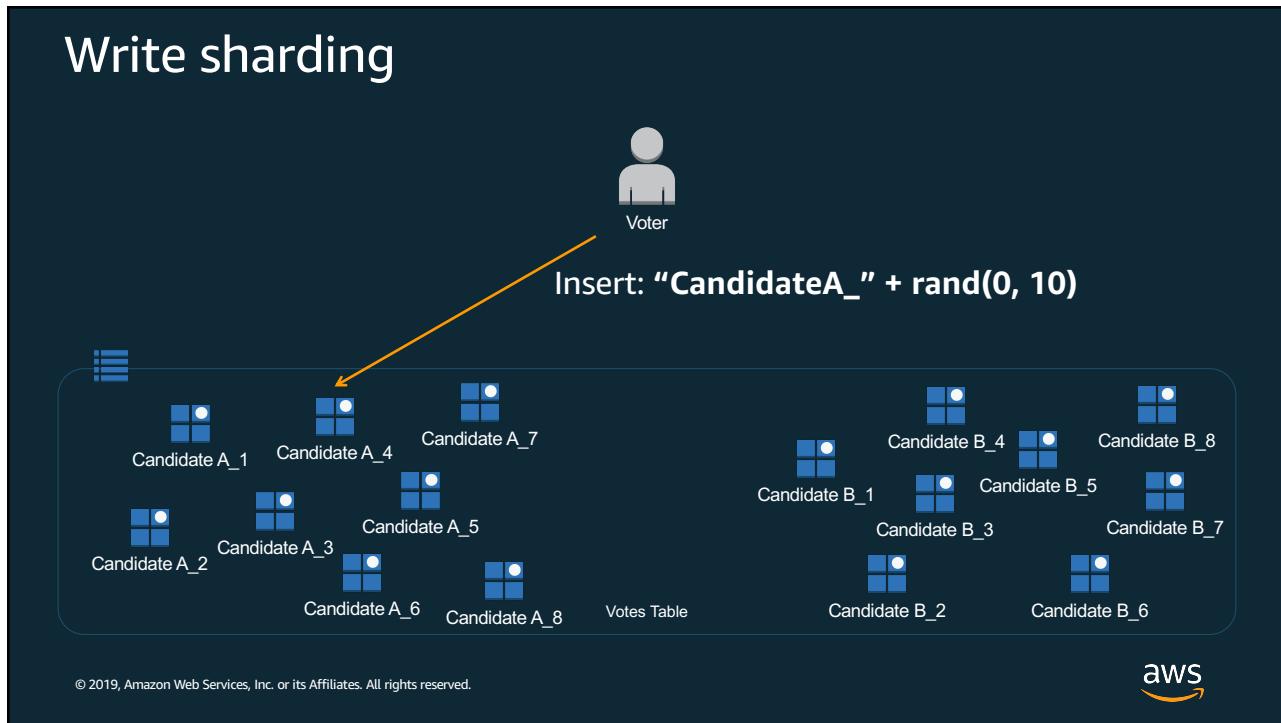
45

## Write sharding

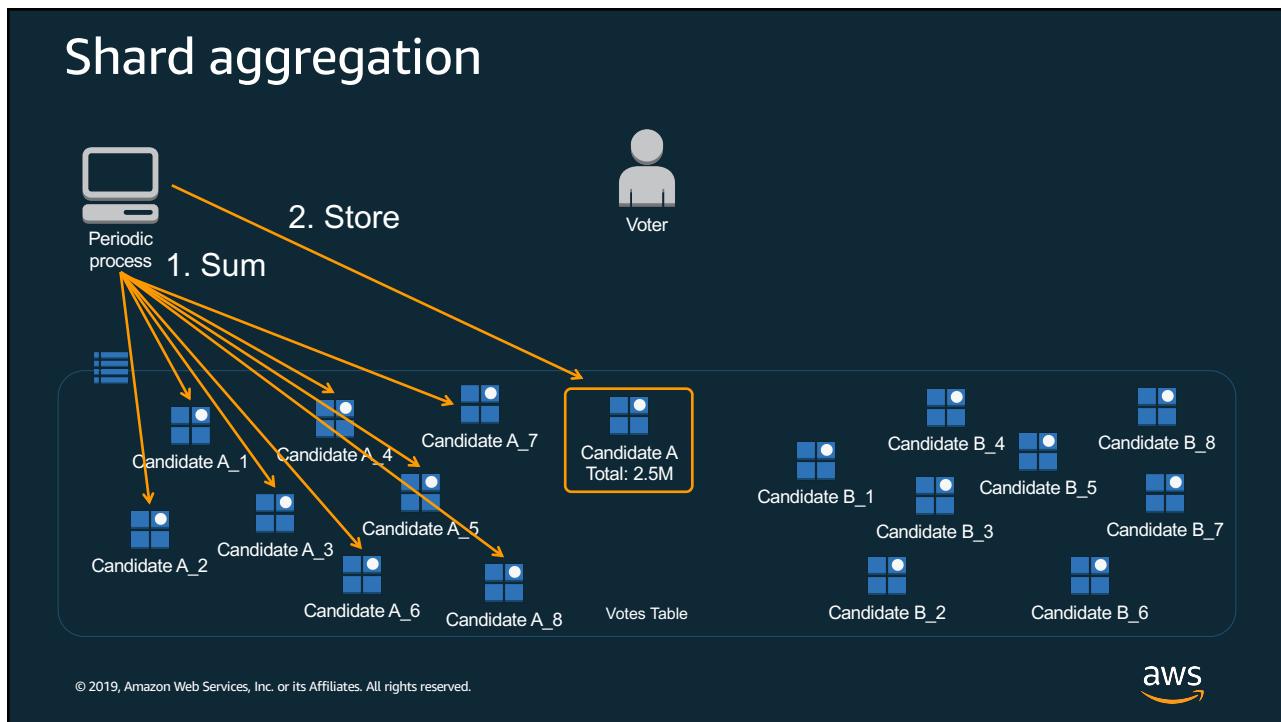


© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

46

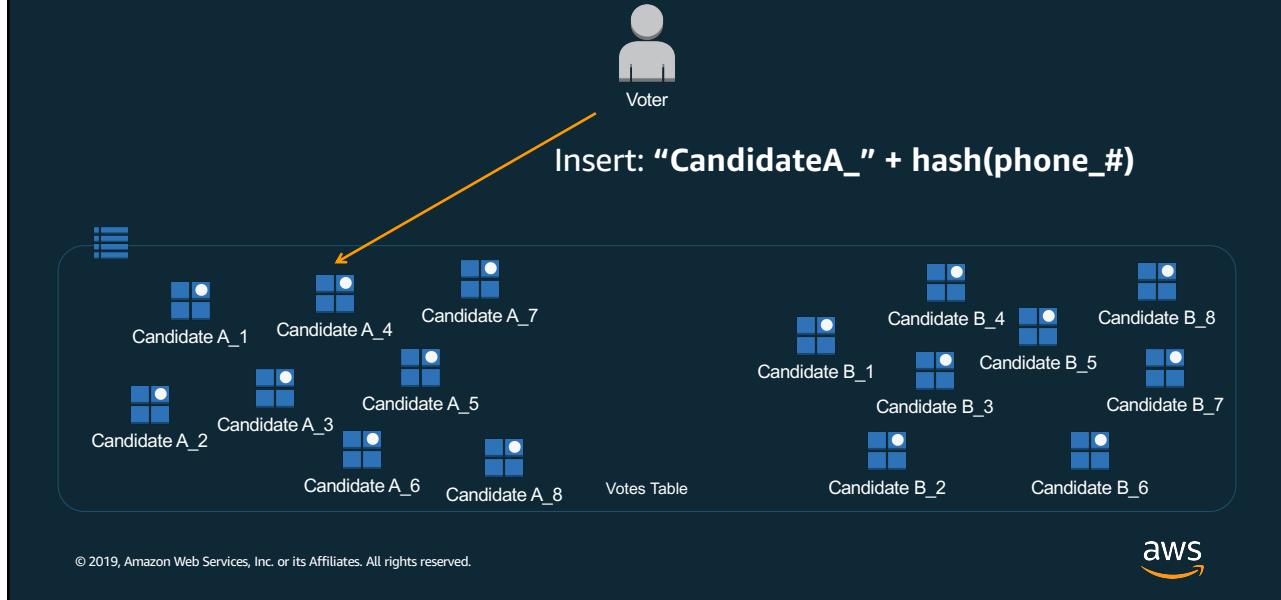


47



48

## Query a single sharded item efficiently



49

## Calculating partition counts (reads)

$$\frac{100 \text{ K} * 0.2 \text{ KB}}{4 \text{ KB}} * 10 / 3000 = \boxed{\sim 17}$$

Items per partition      Average item size      RCU Size      Requests per second      Partition max RCU

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



50

## Calculating partition counts (writes)

$$100 \text{ K} * (\text{itemSize} < 1\text{KB} ? 1\text{KB} : \text{itemSize}) / 1000 = 100$$

Items per second

Average item size

Partition max WCU

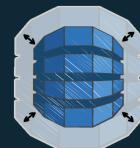
© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



51



### Shard write-heavy partition keys



- Increase throughput with concurrency.
- Evaluate RCU/WCU per key, item size and request rate.



**Important when:** A write workload is not horizontally scalable.

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



52

# Time-based workflows

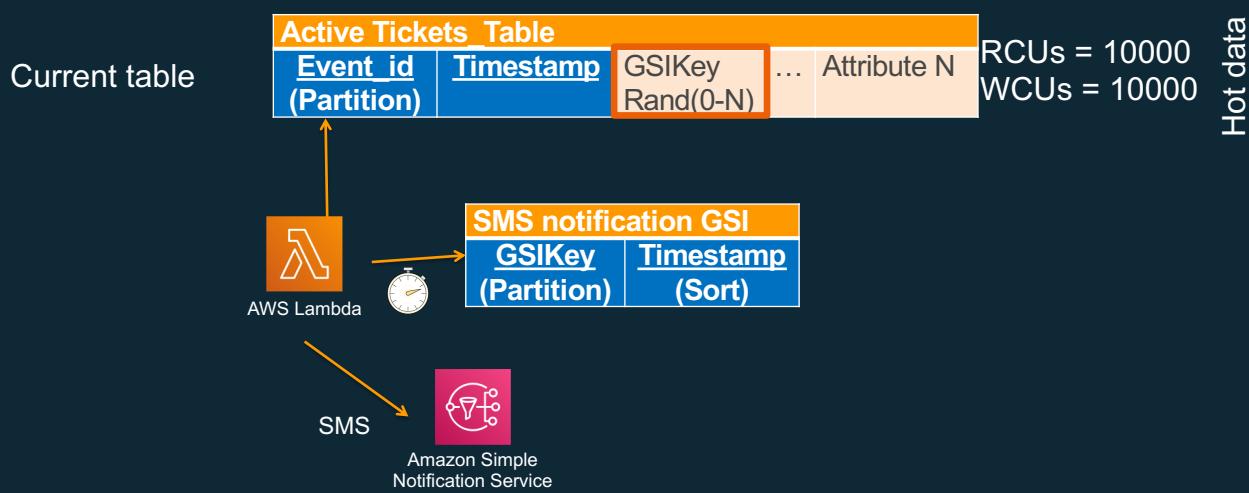
Processing the entire table efficiently

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



53

## Job scheduling



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



54



## Find items efficiently with AWS Lambda

- Write-shard a GSI to selectively query the entire table.
- Use Lambda “stored procedure” to process items.
- Migrate data between tables with AWS Lambda.

**Important when:** You need a massive horizontally scaled index

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



55

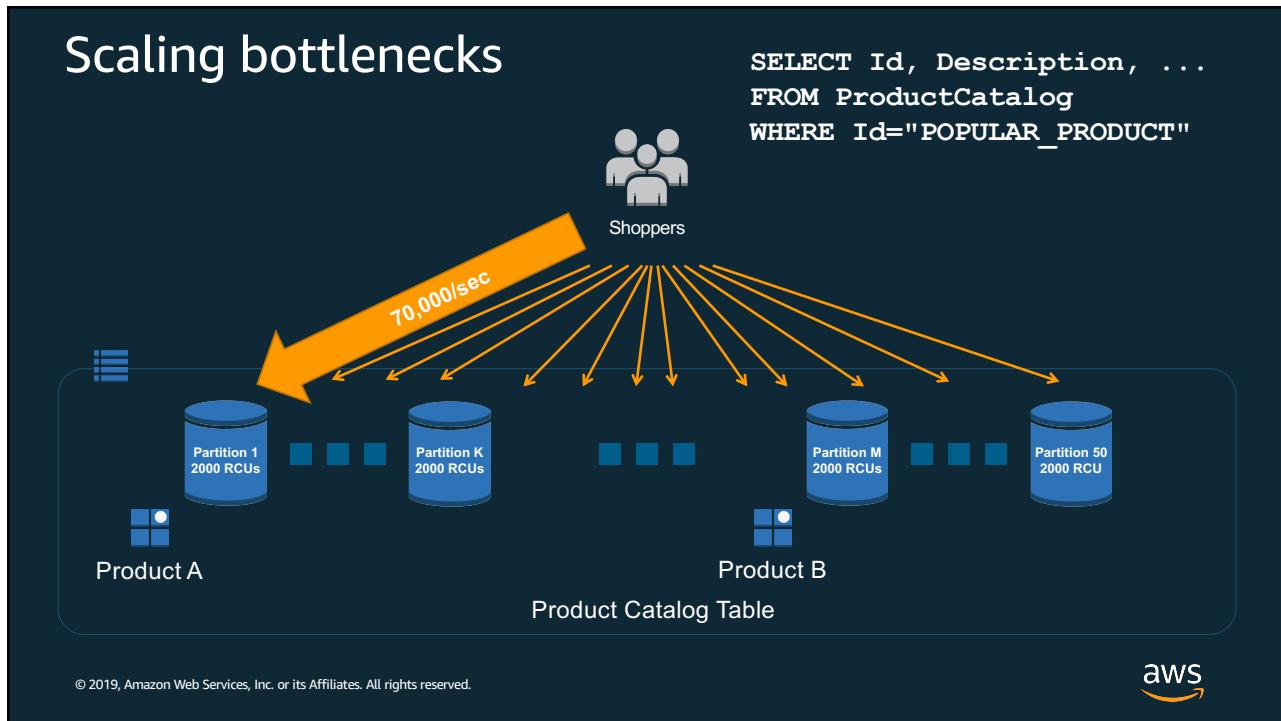
## Product catalog

Popular items (*read heavy*)

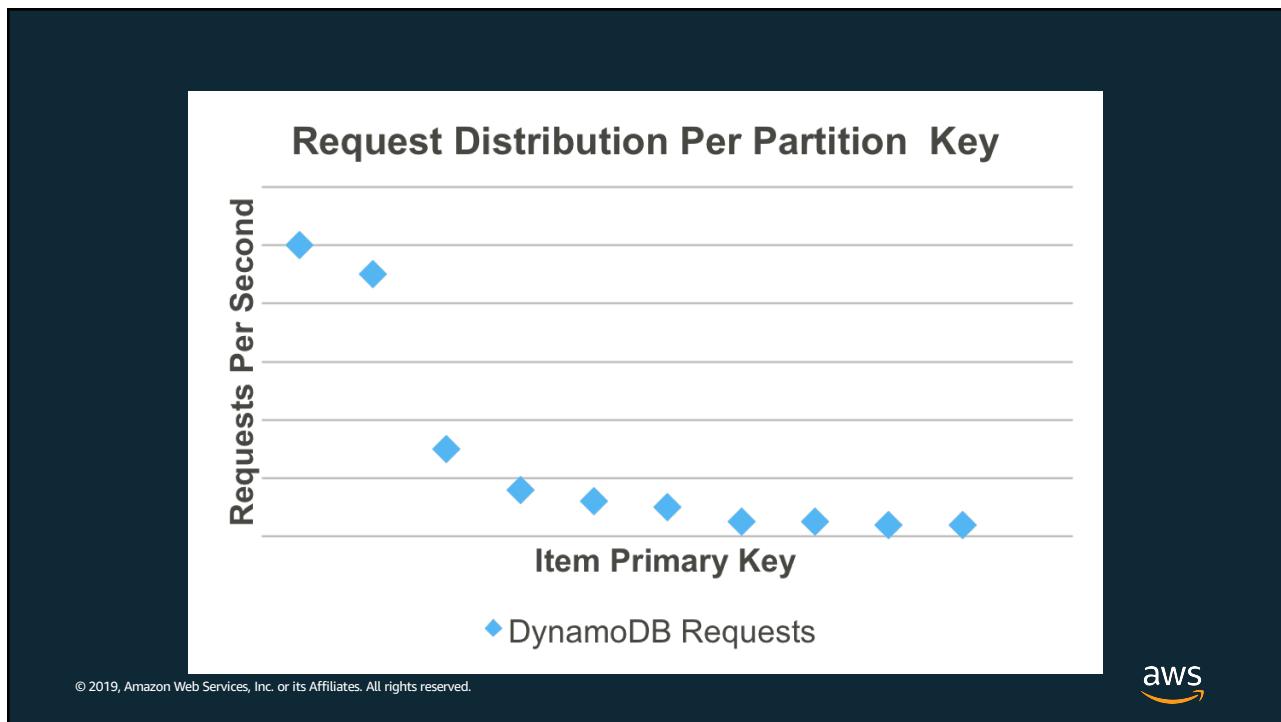
© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



56



57



58

**Your Applications**

DynamoDB Accelerator

DynamoDB

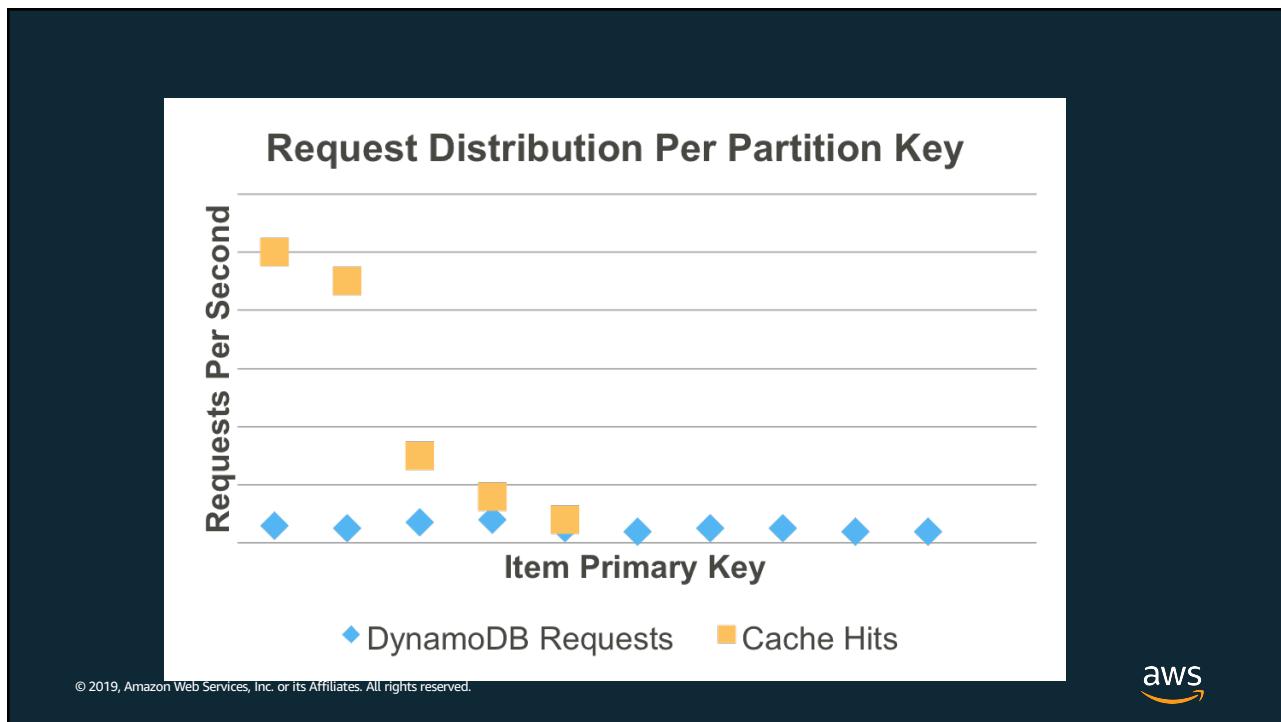
© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

**Features**

- **Fully managed, highly available:** handles all software management, fault tolerant, replication across multi-AZs within a region
- **DynamoDB API compatible:** seamlessly caches DynamoDB API calls, no application re-writes required
- **Write-through:** DAX handles caching for writes
- **Flexible:** configure DAX for one table or many
- **Scalable:** scales-out to any workload with up to 10 read replicas
- **Manageability:** fully integrated AWS service: Amazon CloudWatch, tagging for DynamoDB, AWS Console
- **Security:** Amazon VPC, AWS IAM, AWS CloudTrail, AWS Organizations

aws

59



60

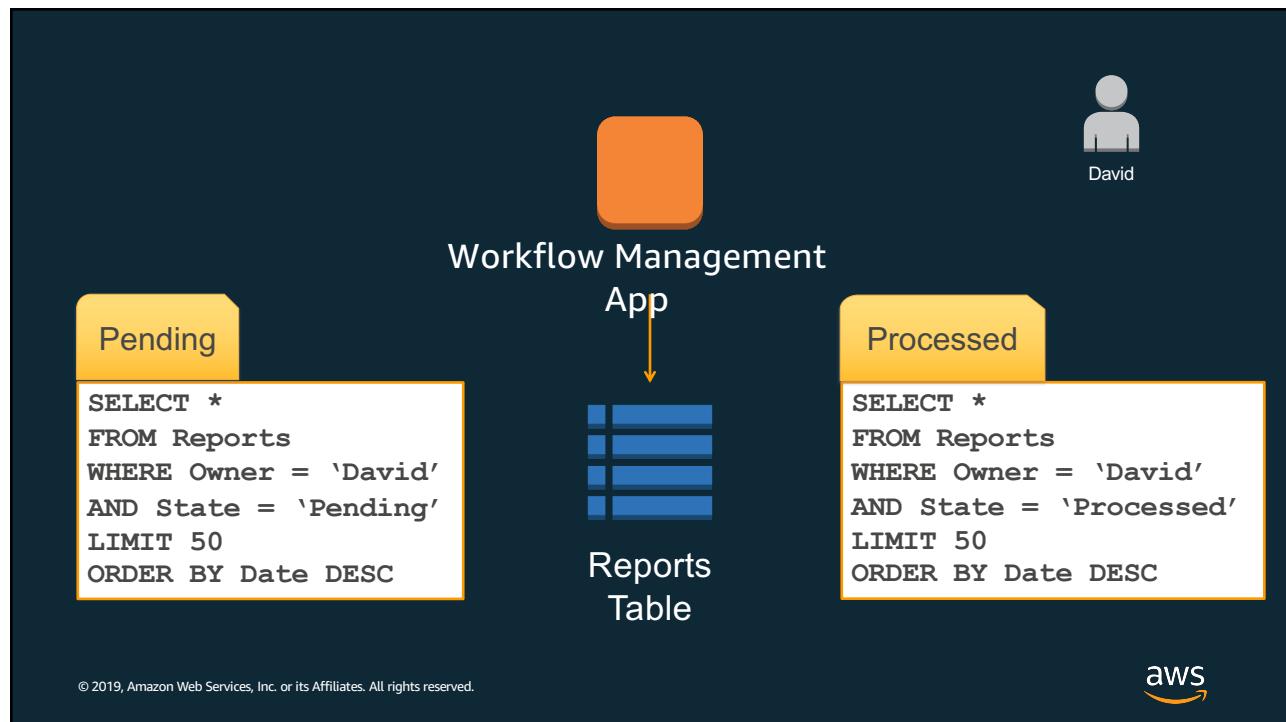
# Vertical Partitioning

Large items  
Filters vs. indexes  
M:N modeling

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



61



62

## Mixed large and small attributes

Partition key      Sort key

Reports table

Owner	StateDate	Document
David	Pending#2014-10-02	...
	... many more Reports for David ...	
David	Processed#2014-10-03	...
Alice	Pending#2014-09-28	...
Alice	Pending#2014-10-01	...

(Many more report items)

Inbox      David

```
SELECT *
FROM Reports
WHERE Recipient='David'
AND Status='Pending'
LIMIT 50
ORDER BY Date DESC
```

50 items × 256 KB each

Large attachments

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

63

## Computing read query cost

$$50 \text{ items} * 256\text{KB item size} * (1 \text{ RCU} / 4 \text{ KB}) * (1 / 2 \text{ conversion ratio}) = 1600 \text{ RCU}$$

Items evaluated by query      Average item size      Conversion ratio      Eventually consistent reads

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

64

## Separate the bulk data

Uniformly distributes large item reads

(50 sequential items at 128 bytes)

1. Query Pending-GSI: 1 RCU
2. BatchGetItem messages: 1600 RCU

(50 separate items at 256 KB)

Pending-GSI

Owner	StateDate	Summary	ReportID
David	Pending#2014-10-02	...	afed
David	Processed#2014-10-03	...	3kf8
Alice	Processed#2014-09-28	...	9d2b
Alice	Processed#2014-10-01	...	ct7r

Reports table

ReportID	Body
9d2b	...
3kf8	...
ct7r	...
afed	...

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.  
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

65

## GIs create selective and efficient queries

Pending

Pending GSI

Processed

Processed GSI

Reports table

David

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

66



## Distribute large items



- Reduce one-to-many item sizes.
- Configure secondary index projections.
- Use GSIs to model M:N relationships.



**Important when:** Querying many large items at once.

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



67

# Relational Data Modeling

Hierarchies & Compound Keys  
Transactions  
Bringing it all together

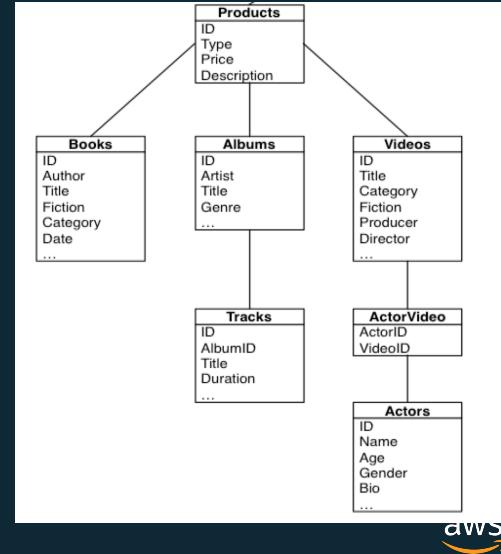
© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



68

## How OLTP apps use data

- Mostly hierarchical structures
- Entity driven workflows
- Data spread across tables
- Requires complex queries
- Primary driver for ACID



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



69

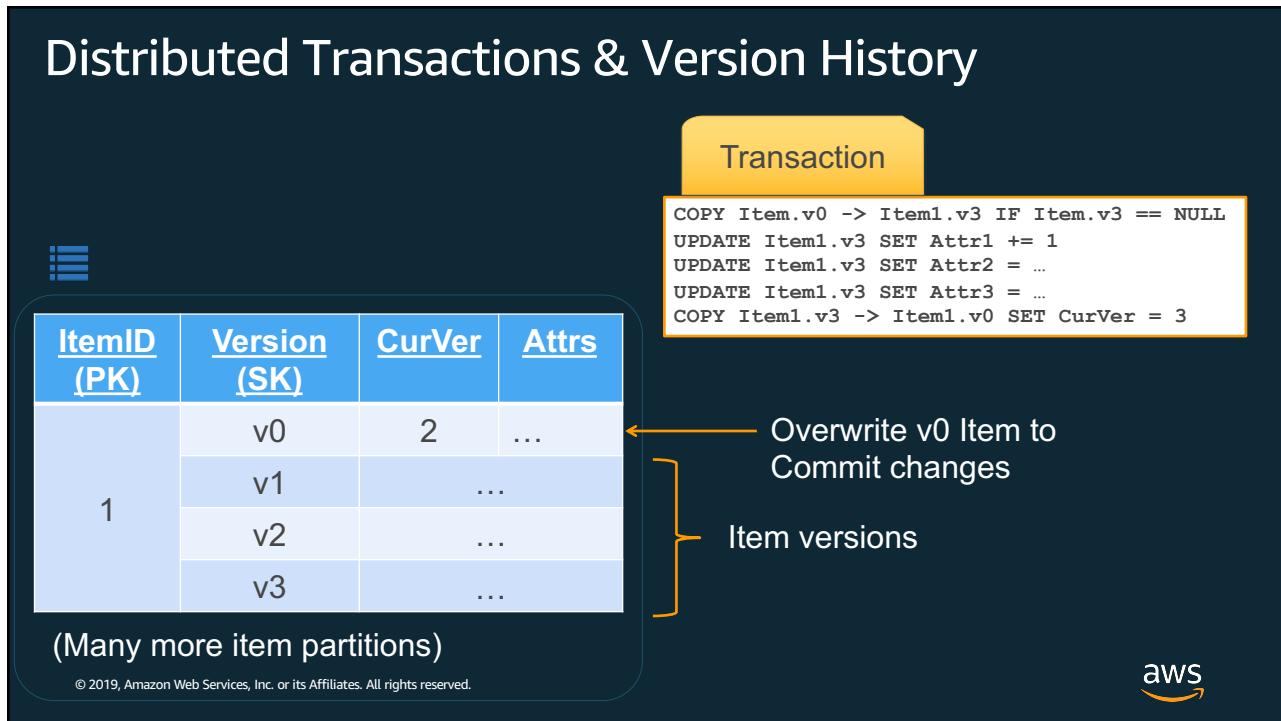
## Multi-version concurrency

Distributed Transactions in NoSQL

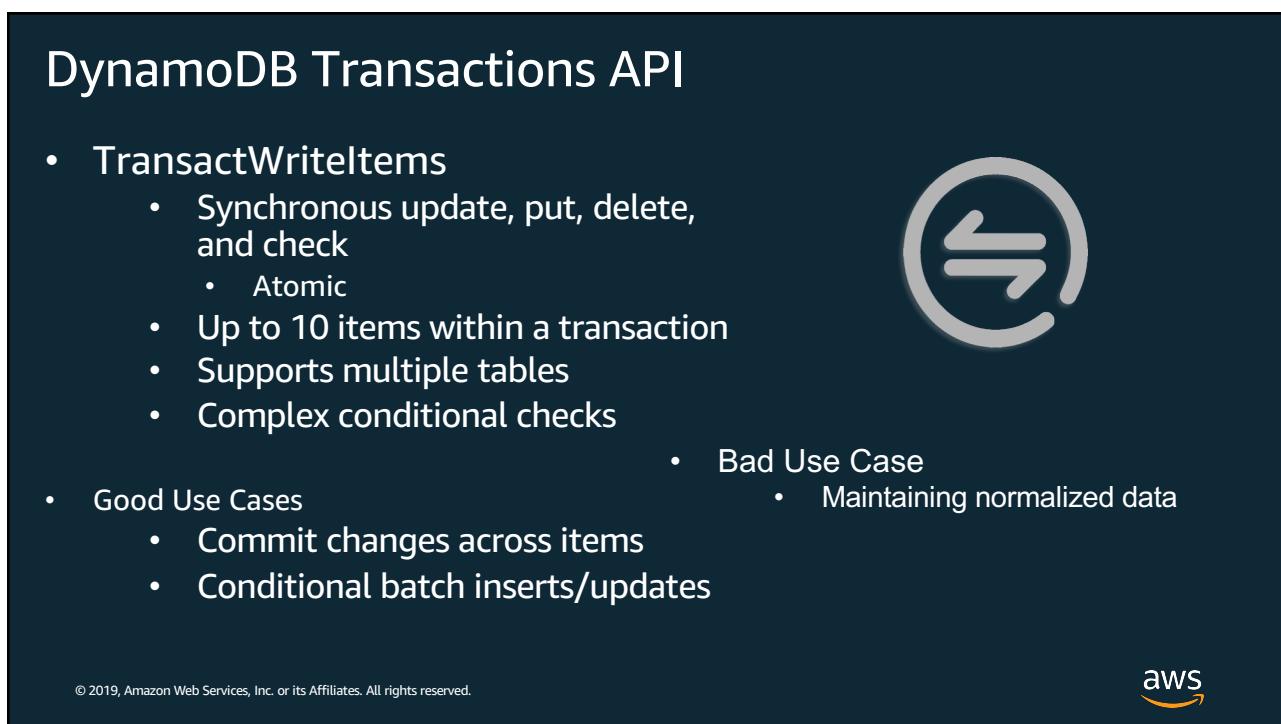
© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



70



71



72

## Game state - Transactions API

*Atomic update of gamer "Hammer57's" Health & Coins*

```
{
  "TransactItems" : [ {
    "Update" : {
      "TableName": "Gamers",
      "Key" : {"GamerTag" : {"S": "Hammer57"}, "Type" : {"S": "Status"}},
      "UpdateExpression" : "Set health = :nhealth",
      "ExpressionAttributeValues":{":nhealth": {"N": "100"}}
    },
    {
      "Update" : {
        "TableName": "Gamers",
        "Key" : {"GamerTag" : {"S": "Hammer57"}, "Type" : {"S": "Assets"}},
        "ConditionExpression" : "coins > :cost",
        "UpdateExpression" : "Set coins = coins - :cost",
        "ExpressionAttributeValues" :{":cost": {"N": "400"}}
      }
    }
  }]
}
```

Gamers					
Primary Key		Attributes			
Gamer Tag	Type	Rank	Level	Points	Tier
Hammer57	Health	87	4050	Elite	
	Status	100%	30		
	Weapon	Taser	87%	50	
	Assets	Coins	1000		
FluffyDuffy	Rank	5	1072	Trainee	
	Status	37	8		
	Assets	Coins	0		

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



73



## Manage transactional workflows

### Distributed Transactions:

- Manage versioning across items with metadata
- Code your app to recognize when updates are in progress
- Implement app layer error handling and recovery logic

### DynamoDB Native Transactions:

- ACID mutations across 10 items

**Important when:** transactional writes are required

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



74

# Hierarchical data

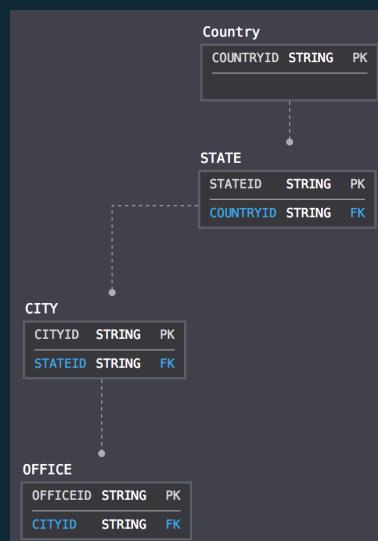
## Composite key modeling

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



75

## Hierarchical data structures as items



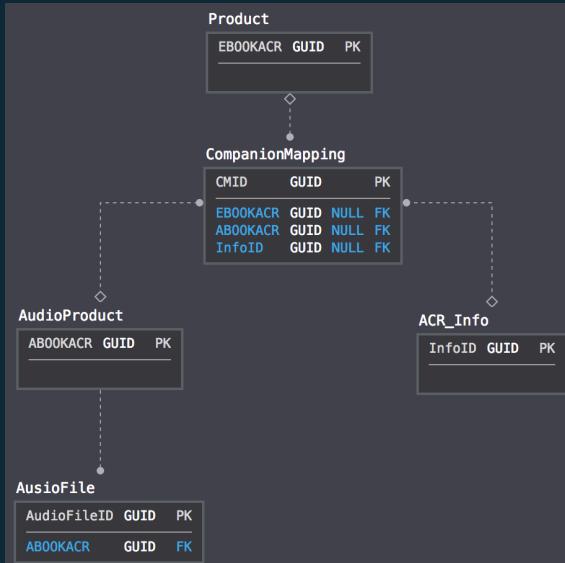
© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

- Use composite sort key to define a hierarchy
- Highly selective queries with sort conditions
- Reduce query complexity

Primary Key		Attributes		
COUNTRY	LOCATION	Address	EmployeeCount	BuildingManager
USA	NY#NYC#JFK11			
	NY#NYC#JFK14			
	WA#SEA#BLACKFOOT			
	WA#SEA#KUMO			
	WA#SEA#MAYDAY			

76

# Audible eBook sync service



- Allows users to save session state for Audible eBooks
- Maintains mappings per user for eBooks and audio products
- Spiky load patterns require significant overprovisioning
- Large number of access patterns



78

# Access patterns

#		USE CASE
1	CompanionMapping	getCompanionMappingsByAsin
2	CompanionMapping	getCompanionMappingsByEbookAndAudobookContentId
3	CompanionMapping	getCompanionMappingsFromCache
4	CompanionMapping	getCompanionMappings
5	CompanionMapping	getCompanionMappingsAvailable
6	AcrInfo	getACRInfo
7	AcrInfo	getACRs
8	AcrInfo	getACRInfos
9	AcrInfo	getACRInfosbySKU
10	AudioProduct	getAudioProductsForACRs
11	AudioProduct	getAudioProducts
12	AudioProduct	deleteAudioProductsMatchingSkuVersions
13	AudioProduct	getChildAudioProductsForSKU
14	Product	getProductInfoByAsins
15	Product	getParentChildDataByParentAsins
16	AudioFile	getAudioFilesForACR
17	AudioFile	getAudioFilesForChildACR
18	AudioFile	getAudioFilesByParentAsinVersionFormat
19	AudioFile	getAudioFiles
20	AudioFile	getAudioFilesForChildAsin

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



79

# Primary table

	Primary Key		Attributes	
	PK	SK (GSI 3)	GSI-1	GSI-2
T A B L E	ABOOKACR1	v0#ABOOKACR1	ABOOK-ASIN1	ABOOK-SKU1
		EBOOKACR1	GSI-1	GSI-2
		ABOOKACR1#TRACK#1	SyncFileAcr	ABOOK-ASIN1
		ABOOKACR1#TRACK#1	GSI-1	GSI-2
		ABOOKACR1#TRACK#2	ABOOK-ASIN1	ABOOK-SKU1
	EBOOKACR1	EBOOKACR1	GSI-1	EBookAsin
			EBOOK-SKU1	ASIN

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



80

# Indexes

G S I 1	Partition Key	Projected Attributes	
	ABOOK-ASIN1	ABOOKACR1	ABOOKACR1-v1 ABOOKACR1#TRACK#1 ABOOKACR1#TRACK#2
G S I 2	SyncFileAcr	ABOOKACR1	MAP-EBOOKACR1
	EBOOK-SKU1	ABOOKACR1	EBOOKACR1

G S I 2	Partition Key	Projected Attributes	
	ABOOK-ASIN1	ABOOKACR1	MAP-EBOOKACR1 ABOOKACR1-v1 ABOOKACR1#TRACK#1 ABOOKACR1#TRACK#2

G S I 3	GSI Partition Key	Projected Attributes	
	V0#ABOOKACR1	ABOOKACR1	ABOOKACR1-v1
	EBOOKACR1	ABOOKACR1	MAP-EBOOKACR1
	ABOOKACR1#TRACK#1	ABOOKACR1	ABOOKACR1#TRACK#1
	ABOOKACR1#TRACK#2	ABOOKACR1	ABOOKACR1#TRACK#2
	EBOOKACR1	ABOOKACR1	EBOOKACR1

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



81

## Query conditions

#	USE CASE	Lookup parameters	INDEX	Key Conditions	Filter Conditions
1	CompanionMapping	getCompanionMappingsByAsin	audiobookAsin/ebookSku	GSI2	GSI-2=ABOOK-ASIN1 None
2	CompanionMapping	getCompanionMappingsByEbookAndAudobookContentId	ebookAcr/sku,version,format or audiobookAcr/asin,version,format	GSI-3 on TargetACR attribute OR PrimaryKey on Table	GSI-3=MAP-EBOOKACR1 version=v and format=f
3	CompanionMapping	getCompanionMappingsFromCache	ebookAcr/sku,version,format or audiobookAcr/asin,version,format	GSI-3 on TargetACR attribute OR PrimaryKey on Table	GSI-3=MAP-EBOOKACR1 version=v and format=f
4	CompanionMapping	getCompanionMappings	syncfileAcr, ebookAcr?, audiobookAcr?	GSI1	GSI-1=SyncFileAcr None
5	CompanionMapping	getCompanionMappingsAvailable	ebookAcr, audiobookAcr	Primary Key on Table	Acr=ABOOKACR1 and TargetACR beginswith "MAP-"
6	AcrInfo	getACRInfo	acr	Primary Key on Table	Acr=ABOOKACR1 and TargetACR beginswith "ABOOKACR1-v"
7	AcrInfo	getACRs	acr / asin,version,format	Primary Key on Table	Acr=ABOOKACR1 version=v and format=f
8	AcrInfo	getACRInfos	acr	Primary Key on table	Acr=ABOOKACR1 and TargetACR beginswith "ABOOKACR1"
9	AcrInfo	getACRInfosbySKU	sku	GSI2	GSI-2=ABOOK-SKU1
10	AudioProduct	getAudioProductsForACRs	acr	Primary Key on table	Acr=ABOOKACR1 and TargetACR beginswith "ABOOKACR1"
11	AudioProduct	getAudioProducts	sku, version, format	GSI2	GSI-2=ABOOK-SKU1 version=v and format=f
12	AudioProduct	deleteAudioProductsMatchingSkuVersion	sku, version	GSI2	GSI-2=ABOOK-SKU1 version=v
13	AudioProduct	getChildAudioProductsForSKU	sku	GSI2	GSI-2=ABOOK-SKU1
14	Product	getProductInfoByAsins	asin	GS1	GSI-1=ABOOK-ASIN1
15	Product	getParentChildDataByParentAsins	asin	GS1	GSI-1=ABOOK-ASIN1
16	AudioFile	getAudioFilesForACR	acr	Primary Key on table	Acr=ABOOKACR1 and TargetACR beginswith "ABOOKACR1#"
17	AudioFile	getAudioFilesForChildACR	acr, parent_asin	Primary Key on table	Acr=ABOOKACR1 version=v and format=f
18	AudioFile	getAudioFilesByParentAsinVersionFormat	parent_asin, version, format	GS1	GSI-1=ABOOK-ASIN1 version=v and format=f
19	AudioFile	getAudioFiles	sku, version, format	GSI2	GSI-2=ABOOK-SKU1 version=v and format=f
20	AudioFile	getAudioFilesForChildAsin	asin, parent_asin, version, format	GS1	GSI-1=ABOOK-ASIN1 version=v and format=f

82

## Recap – what did we cover?

- DynamoDB 101
- Tenants of NoSQL Data Modeling
  - Work backwards from access pattern.
- Common Design Patterns
  - How to build keys to support queries.
  - Indexes, Filters, and Hierarchies:
    - *large items & M:N relationships,*
  - Write Sharding, Distributed Transactions, & Complex Queries
- DynamoDB in the AWS Ecosystem

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



83

# Thank you!

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

